

SYDE 556/750

Simulating Neurobiological Systems
Lecture 13: Conclusion

Chris Eliasmith


Dec 4, 2023

- ▶ Slide design: Andreas Stöckel
- ▶ Content: Terry Stewart, Andreas Stöckel, Chris Eliasmith



UNIVERSITY OF
WATERLOO

FACULTY OF
ENGINEERING



Goal of This Course

Building Large-Scale Brain Models


Why?



Understand how Brains
Work



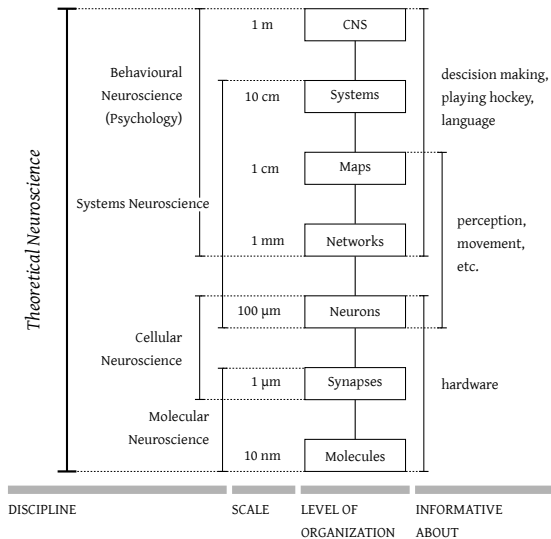
Build Better AI Systems



Program Neuromorphic
Hardware

Image Sources. Left: "A chimpanzee brain at the Science Museum London", from Wikimedia. Centre: "Robot at a campus faire in São Paulo" from Wikimedia. Right: The Braindrop Neuromorphic hardware system, from "Braindrop: A Mixed-Signal Neuromorphic Architecture With a Dynamical Systems-Based Programming Model", Neckar et al., 2019.

Theoretical Neuroscience



The Brain – Some Statistics

- ▶ **Weight:**
2 kg (2% of the body weight)
- ▶ **Power consumption:**
20 W (25% of the body's total power consumption)
- ▶ **Surface area:**
1500 cm² to 2000 cm² (roughly four A4/letter pages of paper)
- ▶ **Number of neurons:**
100 billion (10^{11} , 150 000 mm⁻²)
- ▶ **Number of synapses:**
100 trillion (10^{14} , about 1000 per neuron)

Neuromorphic Hardware

Goal: Brain-inspired hardware; lower power consumption; stream processing


Digital

- ▶ Specialised digital hardware for simulating spiking neural networks
- ▶ Trivial weight-spike multiplication
- ▶ Often asynchronous (no central clock)
- ⊕ Deterministic
- Higher power consumption than analogue

Analogue/Mixed Signal

- ▶ Neuron models in analogue hardware (capacitors, resistors, ...)
- ▶ Digital interconnect and programming (weights, neuron parameters)
- ⊖ Not deterministic
- ⊖ Hard to program
- ⊕ Very low power consumption

Neuromorphic Hardware – SpiNNaker



- ▶ Manchester/Dresden collaboration; HBP
- ▶ Fully digital
- ▶ 18 ARM968 processors @ 180 MHz per chip
- ▶ 1000 current-based LIF neurons per core
- ▶ Toroidal, asynchronous interconnect mesh
- ▶ Up to $\approx 10^9$ neurons in one system
- + Easy to program, (outdated) Nengo interface
- + Public access via HBP
- Not very power efficient (version from 2013)
- High setup times

Neuromorphic Hardware – Loihi




Image Sources. Intel Marketing Material

- ▶ Developed by Intel
- ▶ Digital, fully asynchronous architecture
- ▶ Circuits accelerating individual spiking neurons
- + Extremely low power consumption
- + Nengo Interface
- Proprietary/no low level programming without signing an NDA

Neuromorphic Hardware – BrainDrop






Image Sources. Braindrop: A Mixed-Signal Neuromorphic Architecture With a Dynamical Systems-Based Programming Model, Neekar et al. 2019

- ▶ Research project at Stanford
- ▶ Mixed signal, analogue neurons, synapse arrays
- ▶ Exploits process noise for diverse neural tuning
- ▶ Optimized for NEF networks
- + Extremely low power consumption
- + Nengo Interface
- Availability/Documentation?
- Small networks only

Neuromorphic Hardware – BrainScaleS



- ▶ Research Project in Heidelberg; HBP
- ▶ Mixed signal; above realtime
- ▶ Wafer-scale system; 384×256 neurons

- + Low power consumption
- + Public access via HBP
- Complex; relatively low precision

Review: What Did We Learn? (I)

► The Neural Engineering Framework

- Theory for theoretical neuroscience (**bridging laws**)
- *Principle 1:*
Populations of neurons represent values x
- *Principle 2:*
Connections compute functions f
- *Principle 3:*
Values are states in a dynamical system
- Model of *how* the brain computes. Is it wrong? Of course! But hopefully useful!

Review: What Did We Learn? (II)

► Cognitive Architectures

- *Jackendoff's Challenges:*
How to explain language in neural networks?
- *Vector Symbolic Algebras:*
Compressing symbolic information into vectors,
circular convolution, word embeddings
- *Semantic Pointer Architecture:*
Combination of four ideas – *What* the brain computes.
 - NEF
 - Deep Semantics: compression, decompression
 - Syntax: VSAs (compression, decompression)
 - Architecture: Basal Ganglia/Thalamus/Cortex Loop

Review: What Did We Learn? (III)

► **Methods & Techniques**

- *The Delay Network:*
Efficiently compress past history into a vector; optimal recurrent update rule
- *Spatial Semantic Pointers:*
Bio-inspired representation of continuous spaces; e.g., maps, probabilities
- *Machine Learning:*
Unsupervised, supervised learning, gradient descent,
least squares, nonnegative least squares, delta learning rule (PES)
- *Signal Processing:*
Fourier, Laplace transformation; computing optimal filters
- *Dimensionality Reduction:*
PCA/SVD; function bases; Hebbian learning/Oja learning rule

Review: What Did We Learn? (IV)

► “Meta-Level Skills”

- *Solving Problems by Building a Signal Flow Graph:*
Applications to Hardware design, differentiable computing
- *Programming with Python/Numpy*
- *Building Neural Networks using Nengo:*
Can be applied to neuromorphic hardware (see above), cognitive modeling, machine learning

Summary

► **Party tricks *or* How to impress your mom:**

- *Did you know you can't keep your eyes still in the dark?*

They're controlled by the nuclei prepositus hypoglossi (NPH), part of the brainstem. I built one of those.

- *Did you know people are building hardware that works like the brain?*

It's called neuromorphic hardware, and uses spikes to communicate like the brain. I built some neural networks that can run on that hardware.

- *I know how the world's largest brain model works.*

In fact, I built some of the parts. Let me tell you about the working memory task.

Thank You!